

训练手写数字识别器

计算机视觉第六次作业 | 2101212840 游盈萱

代码

```
1 import torch
2 import torch.nn as nn
3 import torch.utils.data as Data
4 import torchvision
5 import torch.nn.functional as F
6 import numpy as np
7 from torchvision import transforms as tfs
8
9 def train_tf(x):
10     im_aug = tfs.Compose([
11         tfs.Resize(120),
12         tfs.RandomRotation(30),
13         tfs.RandomAffine(degrees=0, translate=(3, 3), scale=None,
14             shear=None, resample=False, fillcolor=0),
15         tfs.ToTensor(),
16     ])
17     x = im_aug(x)
18     return x
19
20 DOWNLOAD_MNIST = True
21
22
23 train_data = torchvision.datasets.MNIST(root='./mnist/', train=True,
24     download=DOWNLOAD_MNIST,
25     transform=train_tf,
26     test_data = torchvision.datasets.MNIST(root='./mnist/', train=False)
27     print(train_data.train_data.shape)
28
29 train_x = torch.unsqueeze(train_data.train_data,
30     dim=1).type(torch.FloatTensor) / 255.
31 train_y = train_data.train_labels
32 print(train_x.shape)
33
34 test_x = torch.unsqueeze(test_data.test_data, dim=1).type(torch.FloatTensor)
35 [:2000] / 255. # Tensor on GPU
36 test_y = test_data.test_labels[:2000]
```

torch.Size([60000, 28, 28])
torch.Size([60000, 1, 28, 28])


```
1 test_x.shape
```

torch.Size([2000, 1, 28, 28])

```

1 import matplotlib.pyplot as plt
2 plt.imshow(test_x[1,0,:,:].numpy(), 'gray')

```



```

1 test_y[:10]

```

tensor([7, 2, 1, 0, 4, 1, 4, 9, 5, 9])

```

1 class NET(nn.Module):
2     def __init__(self):
3         super(NET, self).__init__()
4         self.conv1 = nn.Sequential(nn.Conv2d(in_channels=1, out_channels=16,
5 kernel_size=5, stride=1, padding=2,), nn.ReLU(),
6                                     nn.MaxPool2d(2),)
7         self.dropout1 = nn.Dropout(0.3)
8         self.conv2 = nn.Sequential(nn.Conv2d(16, 32, 5, 1, 2), nn.ReLU(),
9 nn.MaxPool2d(2),)
10        self.dropout2 = nn.Dropout(0.5)
11        self.conv3 = nn.Sequential(nn.Conv2d(32, 64, 4, 1, 2),
12 nn.BatchNorm2d(64), nn.ReLU(), nn.MaxPool2d(2))
13        self.linear1 = nn.Linear(64 * 4 * 4, 10)
14        self.bn = nn.BatchNorm1d(10)
15
16    def forward(self, x):
17        x = self.conv1(x)
18        self.dropout1
19        x = self.conv2(x)
20        self.dropout2
21        x = self.conv3(x)
22        self.dropout2
23        x = x.view(x.size(0), -1)
24        x = self.linear1(x)
25
26        output = self.bn(x)
27        return output
28
29 data_size = 20000
30 batch_size = 32
31 EPOCH = 3
32 LR = 0.002

```

```

32 net = NET()
33
34 optimizer = torch.optim.Adam(net.parameters(), lr=LR)
35 loss_func = nn.CrossEntropyLoss()
36
37
38
39 for epoch in range(EPOCH):
40     random_indx = np.random.permutation(data_size)
41     for batch_i in range(data_size // batch_size):
42         indx = random_indx[batch_i * batch_size:(batch_i + 1) * batch_size]
43
44         b_x = train_x[indx, :]
45         b_y = train_y[indx]
46
47         output = net(b_x)
48         loss = loss_func(output, b_y)
49
50         loss.backward()
51         optimizer.step()
52         optimizer.zero_grad()
53
54         if batch_i % 50 == 0:
55             test_output = net(test_x)
56             pred_y = torch.max(test_output, 1)[1].data.squeeze()
57             # pred_y = torch.max(test_output, 1)[1].data.squeeze()
58             accuracy = torch.sum(pred_y == test_y).type(torch.FloatTensor) /
test_y.size(0)
59             print('Epoch: ', epoch, '| train loss: %.4f' %
loss.data.cpu().numpy(), '| test accuracy: %.3f' % accuracy)
60
61 test_output = net(test_x[:10])
62 pred_y = torch.max(test_output, 1)[1].data.squeeze() # move the computation
in GPU
63
64 print(pred_y, 'prediction number')
65 print(test_y[:10], 'real number')

```

Epoch : 0|trainloss : 2.8338|testaccuracy : 0.403
 Epoch : 0|trainloss : 0.6470|testaccuracy : 0.947
 Epoch : 0|trainloss : 0.5726|testaccuracy : 0.951
 Epoch : 0|trainloss : 0.3787|testaccuracy : 0.965
 Epoch : 0|trainloss : 0.3894|testaccuracy : 0.969
 Epoch : 0|trainloss : 0.4790|testaccuracy : 0.973
 Epoch : 0|trainloss : 0.4183|testaccuracy : 0.977
 Epoch : 0|trainloss : 0.1873|testaccuracy : 0.981
 Epoch : 0|trainloss : 0.2982|testaccuracy : 0.976
 Epoch : 0|trainloss : 0.1809|testaccuracy : 0.978
 Epoch : 0|trainloss : 0.2417|testaccuracy : 0.983
 Epoch : 0|trainloss : 0.3181|testaccuracy : 0.979
 Epoch : 0|trainloss : 0.2491|testaccuracy : 0.980
 Epoch : 1|trainloss : 0.3246|testaccuracy : 0.976
 Epoch : 1|trainloss : 0.2366|testaccuracy : 0.982
 Epoch : 1|trainloss : 0.1431|testaccuracy : 0.982
 Epoch : 1|trainloss : 0.1077|testaccuracy : 0.984
 Epoch : 1|trainloss : 0.2690|testaccuracy : 0.979
 Epoch : 1|trainloss : 0.0928|testaccuracy : 0.982
 Epoch : 1|trainloss : 0.3645|testaccuracy : 0.984
 Epoch : 1|trainloss : 0.3114|testaccuracy : 0.982
 Epoch : 1|trainloss : 0.1202|testaccuracy : 0.985
 Epoch : 1|trainloss : 0.1885|testaccuracy : 0.980
 Epoch : 1|trainloss : 0.1152|testaccuracy : 0.979
 Epoch : 1|trainloss : 0.2172|testaccuracy : 0.985
 Epoch : 1|trainloss : 0.0573|testaccuracy : 0.984
 Epoch : 2|trainloss : 0.1718|testaccuracy : 0.984
 Epoch : 2|trainloss : 0.1892|testaccuracy : 0.984
 Epoch : 2|trainloss : 0.0446|testaccuracy : 0.987
 Epoch : 2|trainloss : 0.0530|testaccuracy : 0.985
 Epoch : 2|trainloss : 0.1286|testaccuracy : 0.989
 Epoch : 2|trainloss : 0.1181|testaccuracy : 0.983
 Epoch : 2|trainloss : 0.0529|testaccuracy : 0.984
 Epoch : 2|trainloss : 0.0915|testaccuracy : 0.984
 Epoch : 2|trainloss : 0.2105|testaccuracy : 0.984
 Epoch : 2|trainloss : 0.0503|testaccuracy : 0.987
 Epoch : 2|trainloss : 0.0714|testaccuracy : 0.987
 Epoch : 2|trainloss : 0.0461|testaccuracy : 0.986
 Epoch : 2|trainloss : 0.0802|testaccuracy : 0.988
 tensor([7, 2, 1, 0, 4, 1, 4, 9, 6, 9])predictionnumber
 tensor([7, 2, 1, 0, 4, 1, 4, 9, 5, 9])realnumber